

ANSON CHAN

SEO, Content & AI Growth Lead

+852 9298 1110 | ai@digitalanson.com | digitalanson.com | Hong Kong — remote-ready, available on-site

English (fluent, writing-grade) · Traditional Chinese (native, writing-grade) · Cantonese (native) · Mandarin (fluent)

PROFESSIONAL SUMMARY

Growth marketer, **8 years**, who has been building **AI-assisted content and campaign production systems in live commercial use since 2022** — two years before the market had a name for it.

I do not run campaigns and hope. **HK\$3M+ cumulative ad spend at 3–10x ROAS**, peaking at **HK\$500K+/month**, every initiative judged on cost per outcome rather than volume. The AI layer is why one operator can hold **10+ tested variants per cycle**, twice-weekly lifecycle sends to 30,000 subscribers, and a monthly reporting pack, without any of the three slipping — **testing cycle time cut ~50%, production time cut ~60%, both measured before and after**.

I write my own copy in **English and Traditional Chinese** to publication standard. And I hold single-variable test discipline, which is the unglamorous reason my numbers are attributable rather than anecdotal.

TWO GAPS, STATED BEFORE ANYTHING ELSE

Your brief says a great candidate brings *"2-3 structured options, clear recommendations"* to every meeting. So here is that treatment applied to my own application.

Gap 1 — I have never worked in B2B SaaS. No ARR, no trial-to-paid funnel, no MQL→SQL→closed-won, no PLG motion. My entire record is B2C and DTC e-commerce, retail, lifestyle and F&B. What transfers is the machine, not the vocabulary: **traffic → activation → conversion → repeat → revenue** is the same funnel your brief describes with different nouns, and I have owned it end to end against a P&L. What does not transfer is knowing, from experience, how a 14-day trial behaves or where mid-market deals stall. **I would be learning that in month one, not bringing it.**

Gap 2 — I am not a technical SEO, and this is the one that matters. Your brief asks for *"flawless technical SEO (indexation, crawlability, Core Web Vitals)"* and *"Schema, entity-driven structures, Information Gain optimization."* My SEO record is channel-level — **30% organic traffic growth** through integrated SEO and paid social — not crawl-budget and log-file level. Claiming otherwise would fall apart inside your first weekly review loop.

What I would propose about Gap 2, rather than just confessing it:

- **Option A (my recommendation):** I own the AI content engine, the editorial standard, GEO/AEO structuring and the funnel maths; the deep technical layer is resourced with a specialist contractor for the first two quarters while I build enough depth to brief and audit it properly. Cheaper than hiring a unicorn, and faster than waiting for one.
- **Option B:** Hire the technical SEO as the lead and me as the AI content-systems half. Two people, clearly split.
- **Option C:** If you need one person who is genuinely flawless on technical SEO *and* AI-native from day one, I am not that person and would rather say so now than in month three.

Everything below is the case for the half I do own.

WHY YOUR BRIEF READS THE WAY IT DOES — AND WHY THAT SUITS ME

Your ad insists that *"every organic initiative must include a target metric, revenue logic, and a clear payback period."* That is not generic rigour. That is a company growing on its own cash rather than on the next round — the sentence a team writes when paid CAC has become too expensive to be the answer.

That is the only way I have ever worked. As a freelancer for eight years, every hour I spent on a workflow had to earn back or I ate the cost. As an in-house manager, my budget was renewed on contribution, not on impressions. I have never had the luxury of an initiative that did not have to pay for itself, so I do not need to be trained into that discipline.

SELECTED EVIDENCE

Built an AI content production system in 2022 — and measured what it saved

At **Baby Central** (Hong Kong's #1 baby e-commerce platform) I put generative AI into client-facing creative and campaign production roughly two years ahead of general adoption. Not a pilot — shipped work.

Built as a **system, not a habit**: prompt libraries per deliverable type, reusable templates, and a named quality gate on every output. Result: **testing cycles cut 50%**, production cost materially reduced, output at **10+ campaign variants**

per cycle while holding single-variable discipline so every result stayed attributable to one change.

This is the "We sell AI, so we must operate with AI" requirement, four years early. And the reason I can quote the savings is that I measured before and after — which is what stops an AI business case from being rejected by whoever holds the budget.

The transferability test I hold myself to

My standard for any AI workflow I build: **someone in their second week produces output at the same standard without asking me**. Undocumented AI capability inside a company is a single point of failure wearing the costume of an advantage — and it is the reason most AI content programmes collapse when the enthusiast leaves.

Relevant to a builder role: the deliverable is the system that survives me, not the output I personally generate.

Agentic tooling, honestly scoped

Claude and Claude Code daily, plus ChatGPT, Gemini, Notion AI, Midjourney, and Zapier / Make-style automation chaining. I build internal tooling with agentic assistance rather than waiting for engineering time — which is the working posture your brief describes.

What I will not overstate: I do not write Python from scratch, and I have not built a programmatic SEO system at scale. I build with agentic tools; I am not a software engineer, and the distinction matters when the brief says *"highly technical."*

Funnel maths, owned against revenue

HK\$500K+/month in Meta and Google spend at **3–7x ROAS** across HK and SEA. **HK\$300K incremental GMV in five months**. Built the **GA4 and Meta Ads Manager dashboards** senior stakeholders reported off monthly — where the real deliverable was getting product, operations and finance to agree what each number meant before anyone was asked to fund the next phase.

Your "Funnel Math" requirement, in the vocabulary I know: ROAS, CAC, CPA, CVR, cohort repeat rate, contribution margin. Different nouns, same arithmetic, and I am comfortable being questioned on any of it.

Rebuilt a failing channel: ROAS 0.3x → 3x in 90 days

A fashion and lifestyle retailer with budget trapped in warm retargeting and a stale creative pool. I moved **70% of spend into cold-audience acquisition** and tested to isolate which *messages* were failing rather than blaming audiences. **3x in 90 days on a majority-acquisition split** — the harder version of that result, because it is net-new demand rather than harvesting what already existed.

"Kill quickly if it doesn't work, scale aggressively if it does" is the JD's phrasing. This is what it looks like when the thing being killed is most of what was already running.

Built organic distribution from zero, with no budget behind it

@aieasyjob — my own Cantonese content brand, **58 films written, directed and cut solo, eight of them in the last month**. One **Facebook Reel reached 76,000 organic views from a page with 77 followers, zero paid support**; a second, built the same way, reached **11,000**, which is what tells you the first was structural rather than luck.

It converts: B2B consulting engagements at HK\$13,800/month — sold off demand I created, to buyers who wanted the AI workflow itself.

One correction before you look, since your brief is about editorial credibility: my channel's lifetime view total is large but roughly 60% of it comes from a single Japanese anime music cover, while the AI and marketing films sit in the hundreds to low thousands. I am deliberately not using the aggregate as evidence. The two Reel figures above are the ones that hold up, and I would rather hand you the honest number than have you find it.

Editorial standards enforced across 200+ third parties

Sourced, briefed and quality-controlled **200+ KOL and influencer partners** across HK and SEA in a trust-sensitive category — including rejecting work before a client saw it. Your brief says *"you enforce high editorial standards that AI cannot replicate on its own."* This is that muscle, built on human contributors, and it transfers directly to governing AI-assisted output at volume.

On leadership

At Baby Central I directly managed a **Marketing Executive and a Designer**, and separately **held two external agencies to account** — one for affiliate marketing, one for SEM — reviewing their plans, checking their numbers, and pushing back when performance didn't hold up. Combined with the 200-partner programme above, that is people management plus vendor management, not one or the other. What I have not done is run a team at SaaS scale with formal weekly review loops — that cadence I would be building, not inheriting.

HOW I WOULD SPEND THE FIRST 90 DAYS

Weeks 1–2 — build the revenue model before touching content. Which organic entry points currently produce trials, which trials convert, at what velocity, and what a marginal visit is worth by segment. Until that model exists, every content decision is a preference rather than an investment. Your brief asks for revenue logic and payback periods on every initiative; this is the layer that makes those numbers real instead of decorative.

Weeks 2–4 — audit the technical floor, with help. Commission an honest technical audit — indexation, crawlability, Core Web Vitals, site architecture, schema coverage. I would rather buy that expertise in week two than pretend to it. The output I need is a prioritised defect list with revenue impact attached, so the fixes get sequenced by value rather than by tidiness.

Weeks 3–8 — build the content engine as a system, not a backlog. Templates by intent type, a research-and-brief pipeline with agentic assistance, and a named quality gate before publication. Target: throughput up several times over with editorial standard **rising**, not holding. The failure mode here is well documented — volume up, quality down, then a Google update erases the year — so the quality gate is not optional garnish.

Weeks 4–12 — attack GEO/AEO deliberately, and measure it. Being the cited source in an AI-generated answer is a different objective from ranking, and most of the market is currently guessing at it. My first move is measurement, not tactics: establish how Wati currently surfaces in Perplexity and ChatGPT Search across the buying-intent question set, so there is a baseline to move. Then structure for it — entity clarity, schema, information gain over restating consensus.

Throughout — shift the narrative up-market, carefully, and bring a kill list to every weekly review. Your brief wants organic aimed at mid-market and enterprise, which is a positioning change before it is a keyword change — enterprise buyers search different questions, in a different register, with procurement in the room. And I would report what I stopped alongside what I started: a growth function with no kill list is not testing, it is accumulating.

PROFESSIONAL EXPERIENCE

Founder & AI Workflow Consultant — @aieasyjob (YouTube · Instagram · Facebook)

2024 – Present | Hong Kong

- Built a Traditional Chinese / Cantonese content brand from zero, solo across three platforms — **58 films to date, eight in the last month** — strategy, script, production, distribution and analytics.
- **76,000-view organic Facebook Reel from a 77-follower page**, zero paid support; second at 11,000.
- Built an AI production pipeline (Claude, Claude Code, ChatGPT, Gemini, Midjourney, Notion AI) with reusable templates and QC gates — **production time cut ~60%**.
- Ran hypothesis-driven tests on hook, message and format against native retention data, compounding findings into a documented playbook.
- Converted organic audience into **B2B consulting engagements at HK\$13,800/month**.

Freelance Digital Marketing Consultant

Mar 2024 – Present | Hong Kong & regional

- **ROAS 0.3x → 3x in 90 days** on a 70% cold-acquisition split for a fashion and lifestyle retailer.
- **+200% conversion rate** through structured single-variable A/B testing at 10+ variants per cycle; peak campaigns above 10x ROAS.
- Monthly client budgets **HK\$30K–200K** at consistent 3–7x ROAS; Klaviyo lifecycle flows driving measurable repeat revenue.

B2C E-Commerce & Digital Marketing Manager

Baby Central Limited (Hong Kong's #1 baby e-commerce platform) / Bloom and Grow Limited

Jan 2022 – Feb 2024 | Hong Kong & SEA

- **Introduced generative AI into commercial creative and campaign production in 2022** — testing cycles cut 50%, production cost materially reduced.
- Owned **HK\$500K+/month** Meta and Google spend at **3–7x ROAS**; **HK\$300K incremental GMV in five months**.
- Grew organic site traffic **30%** through integrated SEO and paid social.
- Owned the **30,000-subscriber Klaviyo lifecycle programme** — twice-weekly campaigns plus welcome, abandoned cart, post-purchase and win-back automation.
- Built **GA4 and Meta Ads Manager dashboards** for monthly senior-stakeholder reporting.
- Built and scaled a **200+ KOL and influencer programme** with briefing standards and output quality control.
- **Directly managed a Marketing Executive and a Designer**, and held **two external agencies to account** (affiliate marketing, SEM) — reviewing plans, checking performance, pushing back when numbers didn't hold.

- Delivered cross-functionally with product, design, development, operations and merchandising on commercial deadlines.

Freelance Digital Marketing Manager — 20+ Retail / Lifestyle / F&B Clients

Nov 2018 – Dec 2021 | Hong Kong

- **HK\$3M+ cumulative ad spend** across 20+ concurrent accounts at 3–7x ROAS.
- Rebuilt funnel architecture, offer structure and messaging register per brand and per purchase cycle; owned pitch, delivery, reporting and renewal solo.

Product Marketing Executive — Mayborn Ltd (Tommee Tippee)

May 2018 – Nov 2018 | Hong Kong

- Social engagement **+15%** via 30+ micro-influencer partnerships generating 3,000+ engagements.
- Contributed to a **5x increase** in organic traffic through integrated digital campaigns.

SKILLS

AI-assisted operations — Claude & Claude Code (daily) · ChatGPT · Gemini · Notion AI · Midjourney · Zapier / Make-style automation chaining · prompt libraries and QC gates per deliverable type · documented, transferable workflows

Content & editorial — English and Traditional Chinese copy written first-hand · hook / message / format testing against retention data · editorial standards enforced across 200+ external contributors · short-form video written, directed and cut solo

Search — SEO delivering 30% organic traffic growth alongside paid social · keyword and intent mapping · landing page messaging and CRO · UTM and pixel governance · *learning: technical SEO depth, programmatic SEO, GEO/AEO measurement*

Funnel & analytics — GA4 · Meta Ads Manager · ROAS, CAC, CPA, CTR, CVR, cohort repeat rate, contribution margin · single-variable A/B testing with kill criteria and winner-scaling rules · dashboard build for non-marketers

Lifecycle — Klaviyo at 30,000 subscribers · welcome, abandoned cart, post-purchase, win-back · segmentation by behaviour rather than demographics

Platforms — Meta Ads · Google Ads (Search, Shopping, Performance Max) · TikTok · Shopify · Klaviyo · Notion (workspace architecture)

CERTIFICATIONS

- **Executive Certificate in Growth Hacking & Data-Informed Marketing — The University of Hong Kong (2019)**
- Advanced Course for Generative AI in Digital Marketing — YMCA College (Dec 2024)
- All-in-One Full-Stack Marketing, No-Code & AI Agent Course (2025 – Present)

EDUCATION

Bachelor of Arts, Travel and Tourism Management — Northumbria University, UK (2005)

Associate Degree, Leisure and Tourism Management — City University of Hong Kong (2004)

LANGUAGES

English — fluent, writing-grade · **Traditional Chinese** — native, writing-grade · Cantonese — native · Mandarin — fluent

AVAILABILITY & PACKAGE

Current: HK\$40,600 per month — self-employed consulting income. Gross, and excluding employer MPF contribution, paid annual leave and medical coverage, so not directly comparable to a permanent package.

Expected: HK\$50,000 – 55,000 per month base, excluding bonus. Morgan McKinley's 2026 Hong Kong reference point for a Digital Marketing Manager is HK\$55,000 per month, and a B2B SaaS growth remit normally carries an uplift on that. **Flexible depending on how the role is levelled and the total package — happy to anchor to your budgeted band rather than guess**, particularly if this Lead role is scoped above the local manager benchmark.

Availability: 2 weeks from agreed terms.

Location: Hong Kong. Remote suits me and I have worked this way for two years; available on-site or hybrid if preferred.